



Tutorial 7: Tree-based Methods

ECO3080: Machine Learning in Business

Instructor: Prof. Qihui Chen
Teaching Assistant: Long Ma

November 15, 2022



- 1 A Simple Classification Tree
- 2 A Simple Regression Tree
- 3 Bagging and Random Forest
- 4 Introduction of Boosting



- The dataset we use in this example is "Carseat".
- We want to predict "Sales" using other features.
- Transform "Sales" into categories:
 - 1 "Sales ≤ 8 ", "High == no"
 - 2 "Sales > 8 ", "High == yes"
- The preprocessing is:

```
#### Get the data we want to use
#### we want to predict the sales of cars by other features
library(ISLR)
Data001 <- Carseats
High <- as.factor(ifelse(Data001$Sales <= 8, "No", "Yes"))
Data002 <- data.frame(Data001, High)
```



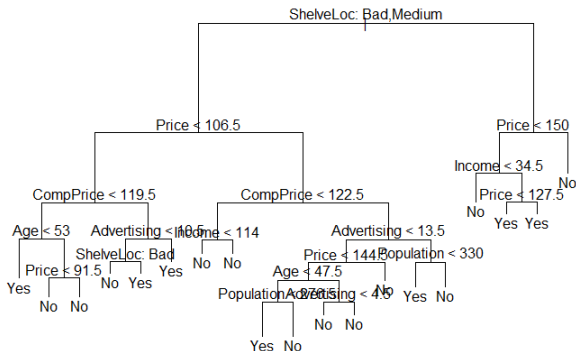
- Separate the dataset into training set and the test set.
- The code is as follows:

```
set.seed(911)
train <- sample(1:nrow(Data002), 200)
Testset <- Data002[-train, ]
High.test <- High[-train]
```

- Then, build up a simple tree on the training set.
- The code is as follows:

```
TrainTree <- tree(High ~ . -sales, Data002, subset = train)
plot(TrainTree)
text(TrainTree, pretty = 0)
```

- The plot is like:





- Use this tree to make prediction on test set:

```
Pred001 <- predict(TrainTree, Testset, type = "class")  
table(Pred001, High.test)
```

- The confusion matrix is:

	High.test	
Pred001	No	Yes
No	87	36
Yes	28	49

- Then, we can calculate a lot of things (sensitivity, lift...). The most straightforward indicator is: $(87+49)/200 = 0.68$



- Then, consider how to prune this tree. CV is feasible. The code is like following:

```
set.seed(1997)
cv.Car <- cv.tree(TrainTree, FUN = prune.misclass)
cv.Car
par(mfrow = c(1, 2))
plot(cv.Car$size, cv.Car$dev, type = "b")
plot(cv.Car$ks, cv.Car$dev, type = "b")
```

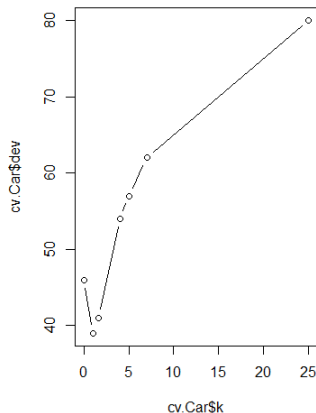
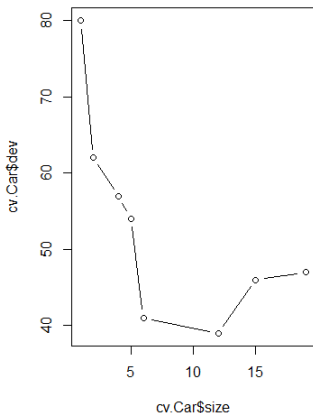
- Thus, we can choose the best number of leaves (From the figures below, 6 or 12 may be good choices. From my perspective, I prefer to choose a simpler tree where size = 6) .

A Simple Classification Tree



香港中文大學(深圳)
The Chinese University of Hong Kong, Shenzhen

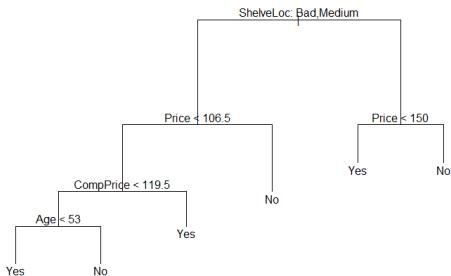
School of Management and Economics, Chinese University of Hong Kong, Shenzhen



- Plug the best "6" into the function "prune.misclass", then we can get a pruned tree.

```
prune.Car <- prune.misclass(TrainTree, best = 6)

par(mfrow = c(1, 1))
plot(prune.Car)
text(prune.Car, pretty = 0)
```





- Use this pruned tree to make prediction on test set.

```
Pred002 <- predict(prune.Car, Testset, type = "class")  
table(Pred002, High.test)
```

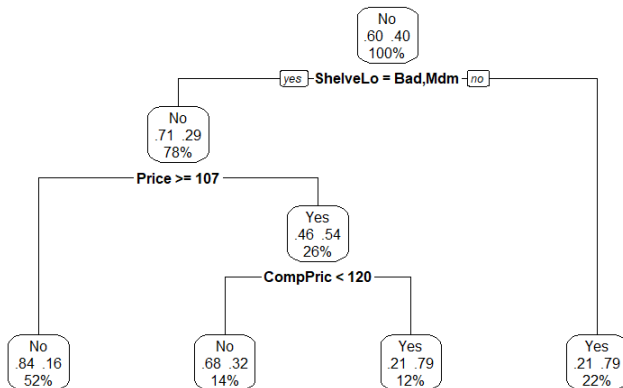
- The confusion matrix is:

	High.test	
Pred002	No	Yes
No	92	38
Yes	23	47

- $(92+47)/200 = 0.695$

- There is another way to build up a tree. "rpart"

Tree 1



■ Same logic:

```
#### Revisit the data set "Boston"
library(MASS)
set.seed(911)
train <- sample(1:nrow(Boston), nrow(Boston)/2)

#### Build up a regression tree
library(tree)
tree.boston <- tree(medv ~ ., data = Boston, subset = train)
summary(tree.boston)

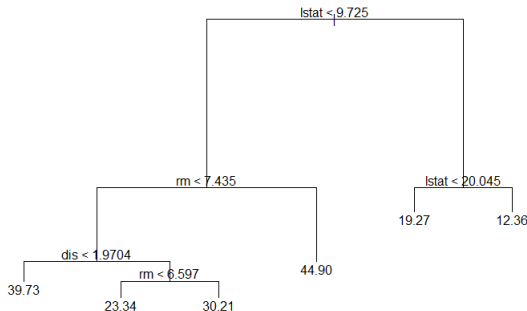
#### Plot this tree
plot(tree.boston)
text(tree.boston, pretty = 0)

#### Use cross validation to get the optimal number of terminal nodes
cv.boston <- cv.tree(tree.boston)
plot(cv.boston$size, cv.boston$dev, type = "b")

prune.boston <- prune.tree(tree.boston, best = 6)
plot(prune.boston)
text(prune.boston, pretty = 0)
```



■ A pruned tree:



■ Prediction:

```
#### Make predictions on test set
yhat1 <- predict(tree.boston, newdata = Boston[-train, ])
yhat2 <- predict(prune.boston, newdata = Boston[-train, ])

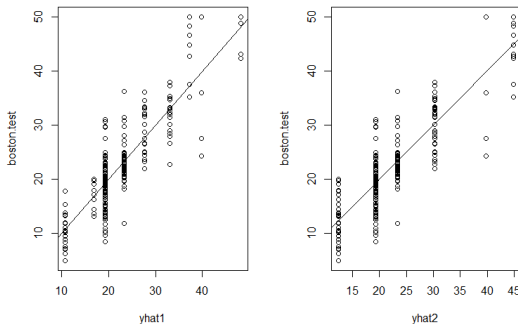
boston.test <- Boston[-train, "medv"]

par(mfrow = c(1, 2))
plot(yhat1, boston.test)
abline(0, 1)

plot(yhat2, boston.test)
abline(0, 1)

mean((yhat1 - boston.test)^2)
mean((yhat2 - boston.test)^2)
```

■ Prediction:



■ $\text{mean}((\text{yhat1} - \text{boston.test})^2) = 20.14976$

■ $\text{mean}((\text{yhat2} - \text{boston.test})^2) = 19.75798$



- You can also use `rpart()` to generate a regression tree;
- Please choose the `"method = "anova"`;
- Also, by using the package `"fancyRpartPlot"`, you are able to make your plots prettier and fancier.

- The code for bagging (from the text book) is:

```
##### 3 Bagging and Random Forest #####
#####
install.packages("randomForest")
library(randomForest)
set.seed(911)
bag.boston <- randomForest(medv ~ ., data = Boston, subset = train, mtry = 13,
                           importance = TRUE) ## why mtry = 13?
####  when mtry = the number of variables, then randomforest == bagging
bag.boston

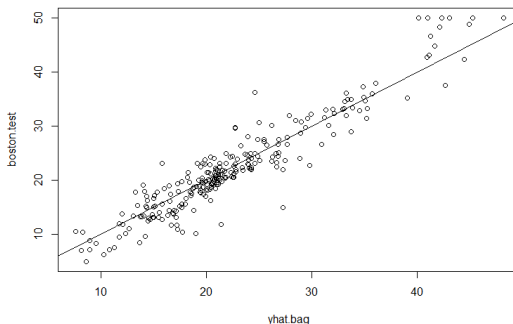
####  make predictions on test set
yhat.bag <- predict(bag.boston, newdata = Boston[-train, ])
par(mfrow = c(1, 1))
plot(yhat.bag, boston.test)
abline(0, 1)
mean((yhat.bag - boston.test)^2)
```

Bagging and Random forest



香港中文大學(深圳)
The Chinese University of Hong Kong, Shenzhen

School of Management and Economics, Chinese University of Hong Kong, Shenzhen



■ The code for RF is:

```
#### change the number of trees
bag.boston <- randomForest(medv ~., data = Boston, subset = train,
                           mtry = 13, ntree = 25)
yhat.bag <- predict(bag.boston, newdata = Boston[-train, ])
mean((yhat.bag - boston.test)^2)

#### randomForest
set.seed(911)
rf.boston <- randomForest(medv ~., data = Boston, subset = train,
                          mtry = 6, importance = TRUE)
yhat.rf <- predict(rf.boston, newdata = Boston[-train, ])
mean((yhat.rf - boston.test)^2)

importance(rf.boston)
varImpPlot(rf.boston)

#### ipred/adabag can also be used in bagging
```

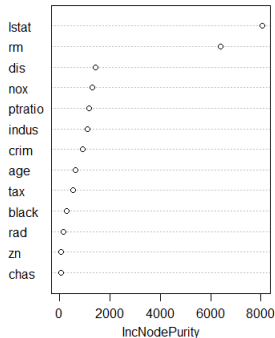
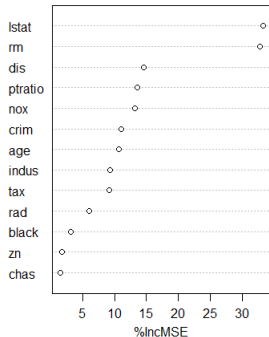
Bagging and Random forest



香港中文大學 (深圳)
The Chinese University of Hong Kong, Shenzhen

School of Management and Economics, Chinese University of Hong Kong, Shenzhen

rf.boston



```
##### 4 Boosting #####
#####
install.packages("gbm")
library(gbm)

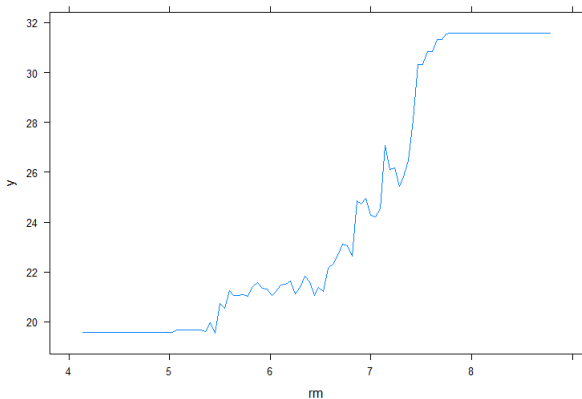
set.seed(911)
boost.boston <- gbm(medv ~ ., data = Boston[train, ], distribution = "gaussian",
                    n.trees = 5000, interaction.depth = 4)
summary(boost.boston)

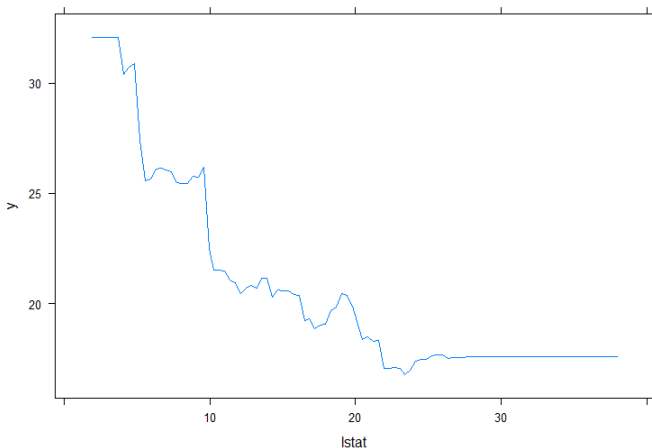
par(mfrow = c(1, 2))
plot(boost.boston, i = "rm")
plot(boost.boston, i = "lstat")

yhat.boost <- predict(boost.boston, newdata = Boston[-train, ], n.trees = 5000)
mean((yhat.boost - boston.test)^2)

boost.boston <- gbm(medv ~., data = Boston[train, ], distribution = "gaussian",
                    n.trees = 5000, interaction.depth = 4, shrinkage = 0.2,
                    verbose = F)
yhat.boost <- predict(boost.boston, newdata = Boston[-train, ], n.trees = 5000)
mean((yhat.boost - boston.test)^2)

#### highly recommend you: adaboost, xgboost .....
#### there are lots of materials on the internet, you can check by yourself
```







- Adaboost, Xgboost